

# Predicting Customer Churn

## Reading the data

The first step is to read the data from the dataset. The dataset consists of various variables such as customer ID, age, gender, account balance, credit score, etc.

In [ ]: `import pandas as pd`

```
df = pd.read_csv("Customer_Churn.csv")
df.head(5)
```

Out [ ]:

	CustomerId	Lastname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts
0	15729836	Robinson	646	Spain	Male	32	1	0.00	2
1	15708610	Costa	690	Germany	Male	44	9	100368.63	2
2	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2
3	15594133	Erskine	697	Spain	Male	62	7	0.00	1
4	15726747	Donaldson	714	France	Male	63	4	138082.16	1

## Dropping unnecessary variables

The next step is to drop unnecessary variables from the dataset that are not relevant to detect customer churn. For instance, the customer ID variable does not provide any useful information in detecting customer churn, and hence it can be dropped.

In [ ]: `df.drop(['CustomerId', 'Lastname'], axis = 1, inplace = True)`  
`df.head(5)`

Out [ ]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	646	Spain	Male	32	1	0.00	2	1	
1	690	Germany	Male	44	9	100368.63	2	0	
2	772	Germany	Male	42	3	75075.31	2	1	
3	697	Spain	Male	62	7	0.00	1	1	
4	714	France	Male	63	4	138082.16	1	0	

## Dummy-encode categorical variables

**Categorical variables such as gender, etc., need to be encoded to numeric values.**

```
In [ ]: df = pd.get_dummies(df, drop_first = True)
df.head(5)
```

```
Out[ ]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	646	32	1	0.00	2	1	0	183289.22
1	690	44	9	100368.63	2	0	0	35342.33
2	772	42	3	75075.31	2	1	0	92888.52
3	697	62	7	0.00	1	1	0	129188.18
4	714	63	4	138082.16	1	0	1	166677.54

## Scaling the data

The next step is to scale the data in the range of (0,1) using the MinMaxScaler() transformation. Scaling the data helps in avoiding bias towards variables with a higher magnitude.

```
In [ ]: ## standardize
from sklearn.preprocessing import StandardScaler, MinMaxScaler

sc = MinMaxScaler()

## fit
X = sc.fit_transform(df.drop('Churned', axis = 1))
```

```
In [ ]: X.min()
```

```
Out[ ]: 0.0
```

```
In [ ]: y = df.Churned
y
```

```
Out[ ]:
```

0	0
1	0
2	1
3	1
4	0
..	
8995	0
8996	0
8997	0
8998	1
8999	0

Name: Churned, Length: 9000, dtype: int64

## Fitting an AutoEncoder architecture

The next step is to fit an AutoEncoder architecture to the scaled data. Remember an AutoEncoder is an unsupervised learning technique that learns the underlying structure of the data by reducing its dimensionality. It consists of an encoder network that maps the input data to a lower-dimensional latent space and a decoder network that reconstructs the original data from the latent space.

```
In [ ]: from tensorflow import keras
        from tensorflow.keras import layers
        import matplotlib.pyplot as plt
```

```
In [ ]: X.shape
```

```
Out[ ]: (9000, 11)
```

```
In [ ]: ## AUTOENCODER
        ## Inputlayer :11
        ## Outputlayer :11, with activation = 'sigmoid' because we used the MinMaxScalar

        autoencoder = keras.Sequential([
            layers.Input([11]), ## number of predictors
            layers.Dense(6, activation = 'relu'),
            layers.Dropout(rate = 0.1),
            layers.Dense(3, activation = 'relu'), ## bottle neck/latent space
            layers.Dropout(rate = 0.1),
            layers.Dense(6, activation = 'relu'),
            layers.Dropout(rate = 0.1),
            layers.Dense(11, activation = 'sigmoid'), ## number of predictors , 0-1 = sigmoid
        ])
```

```
In [ ]: ## Compile
        autoencoder.compile(optimizer = keras.optimizers.Adam(learning_rate = 0.01), loss = 'r
```

```
In [ ]: ## Fit
        ## For this to be an autoencoder all u need to do is fit on X and X
        autoencoder.fit(X,X, epochs = 50, batch_size = 200, validation_split = 0.1)
```

```
Epoch 1/50
41/41 [=====] - 3s 14ms/step - loss: 0.1288 - val_loss: 0.1043
Epoch 2/50
41/41 [=====] - 0s 7ms/step - loss: 0.1044 - val_loss: 0.0884
Epoch 3/50
41/41 [=====] - 0s 5ms/step - loss: 0.0963 - val_loss: 0.0784
Epoch 4/50
41/41 [=====] - 0s 7ms/step - loss: 0.0890 - val_loss: 0.0698
Epoch 5/50
41/41 [=====] - 0s 9ms/step - loss: 0.0842 - val_loss: 0.0653
Epoch 6/50
41/41 [=====] - 0s 8ms/step - loss: 0.0811 - val_loss: 0.0620
Epoch 7/50
41/41 [=====] - 0s 10ms/step - loss: 0.0785 - val_loss: 0.0603
Epoch 8/50
41/41 [=====] - 0s 9ms/step - loss: 0.0766 - val_loss: 0.0584
Epoch 9/50
41/41 [=====] - 0s 8ms/step - loss: 0.0742 - val_loss: 0.0564
Epoch 10/50
41/41 [=====] - 0s 7ms/step - loss: 0.0733 - val_loss: 0.0556
Epoch 11/50
41/41 [=====] - 0s 9ms/step - loss: 0.0723 - val_loss: 0.0545
Epoch 12/50
41/41 [=====] - 0s 11ms/step - loss: 0.0720 - val_loss: 0.0542
Epoch 13/50
41/41 [=====] - 0s 8ms/step - loss: 0.0717 - val_loss: 0.0551
Epoch 14/50
41/41 [=====] - 0s 5ms/step - loss: 0.0708 - val_loss: 0.0523
Epoch 15/50
41/41 [=====] - 0s 5ms/step - loss: 0.0706 - val_loss: 0.0526
Epoch 16/50
41/41 [=====] - 0s 5ms/step - loss: 0.0700 - val_loss: 0.0538
Epoch 17/50
41/41 [=====] - 0s 4ms/step - loss: 0.0699 - val_loss: 0.0522
Epoch 18/50
41/41 [=====] - 0s 5ms/step - loss: 0.0691 - val_loss: 0.0516
Epoch 19/50
41/41 [=====] - 0s 7ms/step - loss: 0.0694 - val_loss: 0.0523
Epoch 20/50
41/41 [=====] - 0s 5ms/step - loss: 0.0687 - val_loss: 0.0510
```

```
Epoch 21/50
41/41 [=====] - 0s 5ms/step - loss: 0.0693 - val_loss: 0.0520
Epoch 22/50
41/41 [=====] - 0s 5ms/step - loss: 0.0692 - val_loss: 0.0505
Epoch 23/50
41/41 [=====] - 0s 6ms/step - loss: 0.0695 - val_loss: 0.0522
Epoch 24/50
41/41 [=====] - 0s 8ms/step - loss: 0.0689 - val_loss: 0.0522
Epoch 25/50
41/41 [=====] - 0s 6ms/step - loss: 0.0688 - val_loss: 0.0521
Epoch 26/50
41/41 [=====] - 0s 5ms/step - loss: 0.0681 - val_loss: 0.0532
Epoch 27/50
41/41 [=====] - 0s 8ms/step - loss: 0.0685 - val_loss: 0.0523
Epoch 28/50
41/41 [=====] - 0s 6ms/step - loss: 0.0680 - val_loss: 0.0531
Epoch 29/50
41/41 [=====] - 0s 4ms/step - loss: 0.0683 - val_loss: 0.0536
Epoch 30/50
41/41 [=====] - 0s 5ms/step - loss: 0.0674 - val_loss: 0.0518
Epoch 31/50
41/41 [=====] - 0s 6ms/step - loss: 0.0680 - val_loss: 0.0527
Epoch 32/50
41/41 [=====] - 0s 5ms/step - loss: 0.0679 - val_loss: 0.0548
Epoch 33/50
41/41 [=====] - 0s 5ms/step - loss: 0.0675 - val_loss: 0.0540
Epoch 34/50
41/41 [=====] - 0s 4ms/step - loss: 0.0677 - val_loss: 0.0507
Epoch 35/50
41/41 [=====] - 0s 5ms/step - loss: 0.0676 - val_loss: 0.0506
Epoch 36/50
41/41 [=====] - 0s 3ms/step - loss: 0.0676 - val_loss: 0.0507
Epoch 37/50
41/41 [=====] - 0s 3ms/step - loss: 0.0684 - val_loss: 0.0488
Epoch 38/50
41/41 [=====] - 0s 3ms/step - loss: 0.0673 - val_loss: 0.0549
Epoch 39/50
41/41 [=====] - 0s 3ms/step - loss: 0.0670 - val_loss: 0.0560
Epoch 40/50
41/41 [=====] - 0s 3ms/step - loss: 0.0675 - val_loss: 0.0552
```

```

Epoch 41/50
41/41 [=====] - 0s 3ms/step - loss: 0.0676 - val_loss: 0.0563
Epoch 42/50
41/41 [=====] - 0s 3ms/step - loss: 0.0676 - val_loss: 0.0548
Epoch 43/50
41/41 [=====] - 0s 3ms/step - loss: 0.0669 - val_loss: 0.0521
Epoch 44/50
41/41 [=====] - 0s 3ms/step - loss: 0.0676 - val_loss: 0.0575
Epoch 45/50
41/41 [=====] - 0s 3ms/step - loss: 0.0673 - val_loss: 0.0549
Epoch 46/50
41/41 [=====] - 0s 3ms/step - loss: 0.0675 - val_loss: 0.0559
Epoch 47/50
41/41 [=====] - 0s 3ms/step - loss: 0.0676 - val_loss: 0.0540
Epoch 48/50
41/41 [=====] - 0s 3ms/step - loss: 0.0672 - val_loss: 0.0578
Epoch 49/50
41/41 [=====] - 0s 3ms/step - loss: 0.0667 - val_loss: 0.0600
Epoch 50/50
41/41 [=====] - 0s 3ms/step - loss: 0.0671 - val_loss: 0.0504
Out[ ]: <keras.callbacks.History at 0x7ff8e520d990>

```

## Determining the 500 most likely to churn customers

**After fitting the AutoEncoder, we can use it to predict the customers who are most likely to churn. We can do this by selecting the customers whose reconstruction error is high. The reconstruction error is the difference between the original input data and the reconstructed data obtained from the AutoEncoder. We can sort the customers based on their reconstruction error and select the 500 customers with the highest error as the most likely to churn.**

```

In [ ]: ## The Reconstruction error
        ## the predictions are the reconstruction of the variables
        predictions = autoencoder.predict(X)

282/282 [=====] - 0s 1ms/step

In [ ]: ## how things were reconstructed
        pd.DataFrame(sc.inverse_transform(predictions))

```

Out[ ]:

	0	1	2	3	4	5	6	7
0	644.085571	37.799175	4.763194	73140.742188	1.526535	9.998091e-01	0.251400	105703.445312
1	651.760010	38.335217	4.972935	81248.335938	1.538330	7.446102e-01	0.322325	100570.234375
2	651.762634	38.335194	4.972907	81247.140625	1.538346	7.446117e-01	0.322788	100571.835938
3	644.124268	37.798306	4.762522	73115.679688	1.526788	9.998106e-01	0.258022	105734.875000
4	652.349365	38.276878	4.866066	55170.718750	1.634177	1.406282e-32	0.997044	104788.875000
...	...	...	...	...	...	...	...	...
8995	644.231567	37.809395	4.767321	73258.757812	1.526786	9.997784e-01	0.252601	105607.625000
8996	661.089661	38.262138	4.875755	77190.625000	1.596354	7.495748e-01	0.998606	106132.828125
8997	659.510437	38.139111	4.829956	74427.531250	1.595359	9.345653e-01	0.998640	107416.062500
8998	651.747559	38.335274	4.973098	81246.960938	1.538267	7.445813e-01	0.320076	100562.015625
8999	649.233826	38.314945	4.968323	74975.125000	1.542635	4.137923e-06	0.233137	100535.351562

9000 rows × 11 columns

MSE Reconstruction Error

In [ ]:


```
import numpy as np
mse = np.mean((X - predictions)**2, axis = 1)
mse
## these are the reconstructed amount
```

Out[ ]:

```
array([0.04287146, 0.13489593, 0.07249676, ..., 0.01971494, 0.12849466,
       0.02695782])
```

In [ ]:

```
plt.hist(np.log(mse));
```



```
In [ ]: df_error = df
df_error['Reconst_Error'] = mse
```

```
In [ ]: # We sort them according to the error so that largest reconstruction are on top
df_error = df_error.sort_values('Reconst_Error', ascending=False)
df_error.head(5)
```

```
Out [ ]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
<b>7980</b>	606	65	10	126306.64	3	0	0	78600
<b>4454</b>	645	68	9	0.00	4	1	1	176350
<b>639</b>	797	55	10	0.00	4	1	1	49410
<b>8144</b>	720	57	1	162082.31	4	0	0	27140
<b>8213</b>	350	39	0	109733.20	2	0	0	123600

**500 customers with the highest error as the most likely to churn.**

```
In [ ]: top500 = df_error.head(500)
```

```
In [ ]: top500.Churned.sum()
```

```
Out [ ]: 185
```

## Verifying the Predictions



**Once we have selected the 500 most likely to churn customers, we need to determine how many of them did actually churn. We can compare our predictions with the actual churn status of the customers to calculate the accuracy of our model.**

```
In [ ]: ## Number of customers total that churned in dataset  
count_ones = df['Churned'].sum()  
count_ones
```

```
Out[ ]: 1822
```

```
In [ ]: 185/1822
```

```
Out[ ]: 0.10153677277716795
```

The 500 observations show us that there were about 10% customers that were detected to be churned.

```
In [ ]: ## Convert the notebook to HTML  
!jupyter nbconvert --to html /content/Project_Part1.ipynb
```

```
[NbConvertApp] Converting notebook /content/Project_Part1.ipynb to html  
[NbConvertApp] Writing 655892 bytes to /content/Project_Part1.html
```